

Introduction to Spatial Databases and SQL

Dylan E. Beaudette

Natural Resources Conservation Service
U.S. Department of Agriculture

—
California Soil Resource Lab (UC Davis)
<http://casoilresource.lawr.ucdavis.edu/>

13 April, 2011




Soils



Talk Outline

- 1 SQL Review
 - History
 - Basic Syntax
 - Table Operations
 - Joins
- 2 Spatial Databases: PostGIS
 - OGC Simple Features
 - GiST Spatial Index
 - PostGIS Metadata Tables
- 3 Spatial SQL
 - Points
 - Lines
 - Polygons
- 4 Data Import/Export
- 5 Examples
 - Soil Texture
 - Soil Water Storage



SQL¹ "Structured Query Language"

History

- developed by IBM in the '70s
- interactive vocabulary for database queries
- most modern systems are built on the 'SQL-92' standard

Modern Uses

- general purpose question *asking* vehicle
- SQL-based interfaces to many types of data: filesystem elements, GIS data, etc.
- often abstracted behind an interface of some kind: i.e. Google, etc.

¹<http://en.wikipedia.org/wiki/SQL>

Flavors of SQL / Portability Issues

Many Vendors / Projects

- client/server: Oracle, MS SQL, Informix, IBM, MySQL, PostgreSQL
- file-based: Access, SQLite, BerkeleyDB

...but all support a subset of the SQL standards

Backwards Compatibility = Not Portable

- standard is vague on actual syntax
- complex & large standard → only subset implemented
- historic deviations from standard preserved

...in most cases the differences are slight

SQL Extensions

Why Bother?

The SQL language is great for simple set operations, but lacks many of the convenient functions found in other languages. Extensions provide the ability to "call" external functions from other common programming languages, entirely within the context of the database.

Some Examples

- "procedural SQL": PL/SQL, SQL PL, PGPLSQL, etc.
- SQL/XML: parsing of XML (extensible markup language documents)
- SQL/R: use of R language commands (numerical algorithms, statistics, etc.)
- SQL/Perl: use of perl language commands libraries (pattern matching, etc.)
- SQL/Python: use of python language commands and libraries

SQL "Structured Query Language"

Syntax Notes

- set-based, declarative computer language
i.e. a program that describes *what* to do, not *how* to do it
- 3-value logic: TRUE, FALSE, NULL
- several language elements:
 - statements: SQL code that has a persistent effect on tables, etc.
 - queries: SQL code that returns data
 - expressions: operations on or tests against a column's contents
 - clauses: logical 'chunks' of statements / queries

```
UPDATE clause [UPDATE country
SET clause  [SET population = population + 1
WHERE clause [WHERE name = 'USA';
```

Expression

Statement

Expression

Predicate

SQL Syntax

Syntax Notes

```
SELECT [ columns ]  
  [ FROM from_item ]  
  [ WHERE condition ]  
  [ GROUP BY expression ]  
  [ HAVING condition ]  
  [ { UNION | INTERSECT | EXCEPT } SELECT [...] ]  
  [ ORDER BY expression [ ASC | DESC ] ]  
  [ LIMIT count ] ;
```

Example Query

```
SELECT column_x, column_y, column_z  
FROM table_x  
WHERE column_x = 'something'  
-- optional  
GROUP BY column_x  
ORDER BY column_x ; -- semi-colon denotes end of SQL statement
```

Operations on a Single Table

The diagram shows two tables, each with 10 rows and 5 columns. In the left table, a red dashed box highlights the first two columns, representing column filtering. In the right table, a red dashed box highlights the first three rows, representing row filtering.

- filtering by column: `SELECT a, b, c, ...`
- filtering by row: `WHERE`
- ordering: `ORDER BY`
- aggregating: `GROUP BY`
- aggregating *then* filtering: `GROUP BY, HAVING`

INSERT/UPDATE/DELETE Statements

INSERT records into a table

```
INSERT INTO chorizon    -- table name
(cokey, hzname, hzdept_r, hzdepb_r) -- record template
VALUES                -- SQL keyword 'here comes the data'
('new_cokey', 'Ap', 0, 10)    -- a new record
```

UPDATE existing records in a table

```
UPDATE chorizon    -- table to modify some records in
SET hzname = 'E'    -- update horizon names to modern conventions
WHERE hzname = 'A2' ;    -- but only the matching historic names
```

DELETE records FROM a table (be careful!)

```
DELETE FROM chorizon    -- table to delete records from
WHERE hzname IS NULL ;    -- records that are missing a horizon name
```

Table Modification Statements³

Altering Table Structure

-- add a new column

ALTER TABLE chorizon **ADD COLUMN** hydrophobicity_index **integer**;

-- now remove the column

ALTER TABLE chorizon **DROP COLUMN** hydrophobicity_index **integer**;

Altering Column Definitions

-- rename a column

ALTER TABLE chorizon **RENAME COLUMN** claytotal_r **TO** clay;

-- change the column's datatype (be careful!)

ALTER TABLE chorizon **ALTER COLUMN** clay **TYPE numeric**;

-- do not allow NULL values in a column

ALTER TABLE chorizon **ALTER COLUMN** clay **SET NOT NULL**;

-- do not allow values over 100%

ALTER TABLE chorizon **ALTER COLUMN** clay **CHECK** (clay <= 100) ;

³<http://www.postgresql.org/docs/8.3/static/sql-altertable.html>

Joins









































Types of Joins

- Cartesian Join: not generally useful, returns all permutation of input rows
- Inner Join: most commonly used, returns rows that occur in **both** tables
 - 1:1 → rows missing from either table omitted
 - 1:many → rows in the left-hand table repeated
 - many:1 → rows in the right-hand table repeated (LU table)
- Left Outer Join: returns all records from the left-hand table, despite missing records in the right-hand table
 - 1:1, 1:many, many:1 → rows missing from right-hand table padded with NULL
- Right Outer Join: same as left outer join, but reversed
 - 1:1, 1:many, many:1 → rows missing from left-hand table padded with NULL

left-hand table

[illegible]

right-hand table



Inner Join

Join map unit data to component data (1:many)

```
SELECT substr(muname, 0, 30) as muname, mapunit.mukey, cokey,  
compname, compct_r  
FROM mapunit JOIN component  
ON mapunit.mukey = component.mukey  
WHERE mapunit.mukey = '464463'  
ORDER BY compct_r DESC;
```

Results

muname	mukey	cokey	compname	compct_r
San Joaquin sandy loam, shall	464443	464443:641360	San Joaquin	85
San Joaquin sandy loam, shall	464443	464443:641362	Exeter	14
San Joaquin sandy loam, shall	464443	464443:641361	Unnamed, ponded	1




Left-Outer Join

Generate a listing of restrictive features for a single map unit

```
SELECT mukey, component.cokey, compname, compct_r, reskind, reshard
FROM
component LEFT JOIN corestrictions
ON component.cokey = corestrictions.cokey
WHERE mukey = '464443'
ORDER BY compct_r DESC;
```

mukey	cokey	compname	compct_r	reskind	reshard
464443	464443:641360	San Joaquin	85	Duripan	Indurated
464443	464443:641362	Exeter	14		
464443	464443:641361	Unnamed, ponded	1		



Joins to Nested Sub-Queries

```
SELECT mukey, mu_area_frac, taxgrtgroup, hd.cokey as id, top, bottom, prop
FROM
(
  -- component weights, in the form of area fractions
  SELECT cd.mukey, cokey, taxgrtgroup, (comp_pct_r::numeric / 100.0) * mu_area as mu_area_frac
  FROM
  (
    -- component keys and percentages
    SELECT mukey, cokey, comp_pct_r, taxgrtgroup
    from component
    where areasymbol = 'ca654'
    and taxgrtgroup is not NULL
  ) as cd
  JOIN
  (
    -- map unit areas by mukey
    SELECT mukey, sum(ST_Area(wkb_geometry)) as mu_area
    from mapunit_poly
    where areasymbol = 'ca654'
    group by mukey
  ) as mu_areas
  on cd.mukey = mu_areas.mukey
) as comp_wts
-- regular join will throw out all components without horizon data
JOIN
(
  -- horizon level data
  SELECT cokey, hzdept_r as top, hzdepb_r as bottom, claytotal_r as prop
  from chorizon
  where om_r is not null
  and areasymbol = 'ca654'
) as hd
on comp_wts.cokey = hd.cokey
ORDER BY mukey, id, top ASC;
```

What is PostGIS?

PostgreSQL

- Relational Database Management System (RDBMS)
- Scalable to n processors, across m computers
- Support for very large data types and tables
- Open Source

PostGIS

- Spatial Extension to PostgreSQL
- Based on C library functions
- OGC Simple Feature Model
- Open Source (compare with \$60K+/CPU for Oracle Spatial)

PostgreSQL



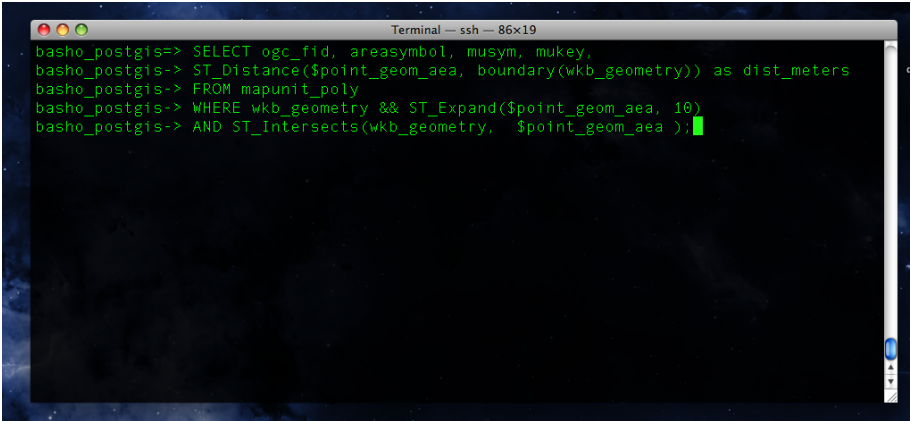
Why Should I Use PostGIS?

- 1 Scales well with massive datasets / file system objects
- 2 Familiar SQL-based manipulation of attribute and spatial features
- 3 Repeatable, transparent work-flow

```
CREATE TABLE dwr.mapunit_dau as
SELECT m_polys.areasymbol, m_polys.mukey, d_polys.dau_id,
ST.Intersection(d_polys.wkb_geometry, m_polys.wkb_geometry) as wkb_geometry
FROM
(
  -- subset map unit polygons to certain survey areas
  -- 6.540 s
  SELECT wkb_geometry, areasymbol, mukey
  FROM mapunit_poly
  -- results in 21682 map unit polygons
  WHERE areasymbol in ('ca653','ca654', 'ca113')
) as m_polys
JOIN
(
  -- subset DAU polygons that overlap with specific survey areas
  -- 512 ms
  SELECT dwr.dau.wkb_geometry, dau.dau3_d_id as dau_id
  FROM dwr.dau JOIN mapunit_bound_poly
  -- results in 58 DAU polygons
  ON mapunit_bound_poly.areasymbol in ('ca653','ca654', 'ca113')
  and dwr.dau.wkb_geometry && mapunit_bound_poly.wkb_geometry
) as d_polys
-- join condition: only those polygons which completely intersect
ON ST.Intersects(d_polys.wkb_geometry, m_polys.wkb_geometry);
```

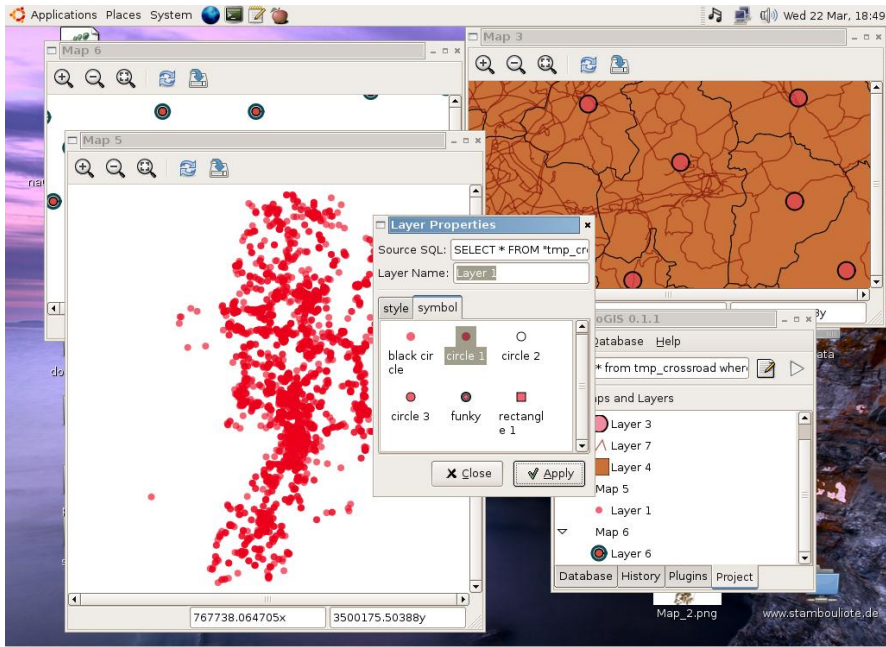
What Does PostGIS Look Like?

psql command-line client

A screenshot of a terminal window titled "Terminal — ssh — 86x19". The terminal has a dark background with green text. It shows a psql command-line session where a user named 'basho_postgis' enters a SQL query. The query selects columns 'ogc_fid', 'areasympbol', 'musym', and 'mukey' from a table named 'mapunit_poly'. It also includes a subquery for 'ST_Distance(\$point_geom_aea, boundary(wkb_geometry))' aliased as 'dist_meters'. The query is filtered by 'WHERE wkb_geometry && ST_Expand(\$point_geom_aea, 10)' and 'AND ST_Intersects(wkb_geometry, \$point_geom_aea);'. A green cursor is visible at the end of the last line of the query.

```
basho_postgis=> SELECT ogc_fid, areasympbol, musym, mukey,  
basho_postgis-> ST_Distance($point_geom_aea, boundary(wkb_geometry)) as dist_meters  
basho_postgis-> FROM mapunit_poly  
basho_postgis-> WHERE wkb_geometry && ST_Expand($point_geom_aea, 10)  
basho_postgis-> AND ST_Intersects(wkb_geometry, $point_geom_aea);
```

What Does PostGIS Look Like?



OGC Simple Feature Model

The Standard

- “Simple Features” - not a topological format (i.e. GRASS or Coverage)
- 2D and 3D geometric primitives, without self-intersection
- <http://www.opengeospatial.org/standards/sfs>

Basic Geometries

- POINT ($x\ y$)
- LINESTRING ($x_1\ y_1, x_2\ y_2, x_3\ y_3, \dots, x_n\ y_n$)
- POLYGON ($x_1\ y_1, x_2\ y_2, x_3\ y_3, \dots, x_{n-1}\ y_{n-1}, x_1\ y_1$);

“Multi-” Geometries

- MULTIPOINT ($(POINT_1), \dots, (POINT_n)$)
- MULTILINESTRING ($(LINESTRING_1), \dots, (LINESTRING_n)$)
- MULTIPOLYGON ($(POLYGON_1), \dots, (POLYGON_n)$) ;

ANSI SQL Examples

- **SELECTION** `SELECT a from b where c = d`
- **sorting** `SELECT a from b order by a desc`
- **join** `SELECT t1.a, t2.b from t1 join t2 on ...`
- **aggregation** `SELECT sum(a)from b group by a`

OGC “Spatial” SQL Examples

- **feature extraction** `SELECT ST_X(point_geom), ST_Y(point_geom)from ...`
- **feature extraction** `SELECT PointN(geom)from ...`
- **spatial join** `SELECT * from t1 join t2 on ST_Distance(t1.geom, t2.geom)< 100 ...`
- **feature manipulation** `SELECT ST_Transform(geom, SRID)from ...`
- **feature analysis** `SELECT ST_Buffer(geom, distance)from ...`
- **GIS overlay functions** `SELECT ST_Intersection(geom_1, geom_2)from ...`

Spatial SQL: Point Geometry

Construction

- **SELECT** ST_MakePoint(x,y)
- **SELECT** ST_Centroid(polygon_geom)
- **SELECT** ST_PointOnSurface(polygon_geom)

Measurement

- **SELECT** ST_X(geom), ST_Y(geom)
- **SELECT** ST_Distance(geom1, geom2)
- **SELECT** ST_Distance_Sphere(geom1, geom2)

Geometric Operation

- **SELECT** ST_Buffer(geom, distance)
- **SELECT** ST_Expand(geom, distance)
- **SELECT** ST_Touches(geom1, geom2)

Spatial SQL: Line Geometry

Construction

- **SELECT** ST_MakeLine(geometry collection)
- **SELECT** ST_MakeLine(geom1, geom2)

Measurement

- **SELECT** ST_Length(geom)
- **SELECT** ST_Length_Spheroid(geom, spheroid)
- **SELECT** ST_Length3d(geom)

Geometric Operation

- **SELECT** ST_Crosses(geom1, geom2)
- **SELECT** ST_Overlaps(geom, distance)
- **SELECT** ST_Intersection(geom1, geom2)

Spatial SQL: Polygon Geometry

Construction

- **SELECT** ST_ConvexHull(geometry collection)
- **SELECT** ST_BuildArea(line_geom)

Measurement

- **SELECT** ST_Area(geom)
- **SELECT** ST_Perimeter(geom)

Geometric Operation

- **SELECT** ST_Intersection(geom, geom)
- **SELECT** ST_Intersects(geom, geom)
- **SELECT** ST_Contains(geom1, geom2)

Getting Data into and out of PostGIS

OGR tools: recall order of OGR data sources: *output input*

■ Import

```
ogr2ogr -f "PostgreSQL" \  
PG:'dbname=ssurgo_combined user=xxxx password=xxxx host=host' input_file.shp
```

■ Export

```
ogr2ogr output_file.shp \  
PG:'dbname=ssurgo_combined user=xxxx password=xxxx host=postgis.server.edu' tablename
```

→ Note the OGR syntax for specifying a PostGIS DSN.

PostGIS Loader/Dumper

■ Import

```
shp2pgsql -s SRID -c -g wkb_geometry -I shapefile.shp schema.table \  
| psql -U username -h host database
```

■ Export

```
pgsql2shp -f shapefile.shp -h host -u username -P password -k -g wkb_geometry \  
database schema.table
```

→ See the manual page for `pgsql2shp` for a complete list of arguments.

Getting Data into and out of PostGIS

Text Files

- CSV format, from within the psql client

```
\copy tablename TO 'filename.csv' CSV HEADER
```

- CSV format, via psql client

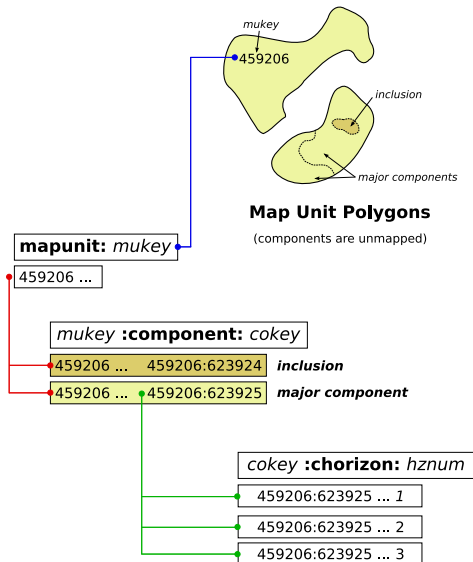
```
echo "select column_list from table_list" | psql --tuples --no-align -F "," database > file.csv
```

- Tabular data to HTML format, via psql client See output below:

```
echo "select column_list from table_list" | psql --html database > file.html
```

```
area compname
132472.230854819 Hilmar variant
322819.967391312 Oneil
362729.418301135 Carranza
431948.171760353 Tuff rockland
448784.927049035 Gravel pits
500763.225267798 Snelling variant
518860.954990617 Foster
571640.132661382 Alamo
648973.748756059 Toomes
924327.631201791 Dumps
```

SSURGO Table Diagram



Aggregation of SSURGO Geometry

Query: extract SSURGO geom. from arbitrary bbox, compute area weights

-- define a transformed bounding box for later use

```
\set bbox ST_Transform(  
ST_SetSRID(  
ST_MakeBox2D( ST_MakePoint(-122.25, 39.28), ST_MakePoint(-122.20, 39.30) )  
, 4326)  
, 9001)
```

-- select map unit keys, map unit symbols, and computed areas for the intersecting polygons

```
SELECT mukey,  
sum(ST_Area(ST_Intersection(wkb_geometry, :bbox))) / ST_Area(:bbox) as mu_area_wt  
from mapunit_poly  
WHERE ST_Intersects(wkb_geometry, :bbox)  
GROUP BY mukey;
```

mukey	mu_area_wt
461544	0.562595368617999
461571	0.347993963186697
461595	0.0748614412770969
461667	0.0145492269180839

Soil Texture Example

compute several weighted means of sand, silt, clay

-- join with polygons, and compute areas weights

```
SELECT mapunit_poly.mukey,  
sum(ST_Area(wkb_geometry)) /  
(SELECT ST_Area(wkb_geometry) FROM mapunit_bound_poly WHERE areasymbol = 'ca113') AS area_wt,  
max(sand) as sand, max(silt) as silt, max(clay) as clay -- fake aggregate functions applied
```

FROM

mapunit_poly

JOIN

(

-- compute component percent weighted mean

SELECT mukey,

sum(comp_pct_r * sand) / sum(comp_pct_r) AS sand,

sum(comp_pct_r * silt) / sum(comp_pct_r) AS silt,

sum(comp_pct_r * clay) / sum(comp_pct_r) AS clay

FROM

component

JOIN

(

-- compute hz thickness weighted mean

SELECT cokey,

sum((hzdepb_r - hzdept_r) * sandtotal_r) / sum(hzdepb_r - hzdept_r) AS sand,

sum((hzdepb_r - hzdept_r) * silttotal_r) / sum(hzdepb_r - hzdept_r) AS silt,

sum((hzdepb_r - hzdept_r) * claytotal_r) / sum(hzdepb_r - hzdept_r) AS clay

FROM horizon

GROUP BY cokey

) AS co_agg

ON component.cokey = co_agg.cokey

GROUP BY component.mukey


) AS mu_agg

ON mapunit_poly.mukey = mu_agg.mukey

GROUP BY mapunit_poly.mukey;

Soil Texture Example (cont.)

Yolo County Soil Textures



Soil Texture Example (cont.)

```
# load some libs:
library(plotrix)

# read in the data
x <- read.csv('yolo_texture.csv')

# simple soil texture, with symbol size weighted by area weight
soil.texture(x[,3:5], cex=sqrt(50*x$area_wt), pch=16, col.symbol=rgb(65,105,225, alpha=100, max=255),
show.lines=T, show.names=T, col.lines='black', col.names='black', main='Yolo County Soil Textures')

triax.points(cbind(weighted.mean(x$sand, x$area_wt), weighted.mean(x$silt, x$area_wt), weighted.mean(x
  $clay, x$area_wt)),
col.symbols='orange', pch=16, cex=2)
```

Soil Water Storage Computation

```
SELECT mukey, compname, compct_r, a.* FROM component
JOIN
(
  SELECT cokey, sum( (hzdepb_r - hzdept_r) * awc_r) AS component_whc, sum((hzdepb_r -
    hzdept_r)) AS depth
  FROM chorizon WHERE areasymbol = 'ca113'
  GROUP BY cokey
) AS a
ON component.cokey = a.cokey
WHERE component.areasymbol = 'ca113'
ORDER BY mukey ;
```

mukey	compname	compct_r	cokey	component_whc	depth
459204	Gravel pits	100	459204:659832	3.04	152
459206	Arbuckle	70	459206:623924	16.46	152
459206	Arbuckle	15	459206:1128332	16.46	152
459207	Arbuckle	85	459207:623932	16.46	152
459208	Balcom	85	459208:623933	9.69	61
459209	Balcom	85	459209:623937	17.25	104
459210	Balcom	85	459210:623942	17.25	104
459211	Balcom	85	459211:623949	9.44	61
459212	Balcom	45	459212:623950	13.93	86

Soil Water Storage Computation

```
SELECT mukey,
-- note that weights from non-soil components must be removed
-- otherwise, weighted mean values will be too low
SUM(compct_r * component_whc) / SUM(compct_r) AS compct_weighted_whc_cm
FROM component
JOIN
(
  SELECT cokey, sum( (hzdepb_r - hzdept_r) * awc_r) AS component_whc,
  sum((hzdepb_r - hzdept_r)) AS depth
  FROM chorizon
  WHERE areasymbol = 'ca113'
  GROUP BY cokey
) AS a
USING (cokey)
WHERE component.areasymbol = 'ca113'
-- filter out components that are missing soils data
AND a.component_whc IS NOT NULL
GROUP BY mukey ;
```

```
mukey | compct_weighted_whc_cm
-----+-----
459225 | 10
459226 | 10
459227 | 11
...
```

Soil Water Storage Computation

```
-- create the new table with both geometry and attributes
CREATE TABLE yolo.whc AS
SELECT ogc.fid, wkb.geometry AS wkb_geometry, b.mukey, b.compct_weighted_whc_cm
FROM mapunit.poly
-- use LEFT JOIN to include non-soil polygons in the result set
-- alternatively use JOIN to ignore non-soil polygons
LEFT JOIN
(
  SELECT mukey,
  -- note that weights from non-soil components must be removed
  -- otherwise, weighted mean values will be too low
  SUM(compct_r * component_whc) / SUM(compct_r) AS compct_weighted_whc_cm
  FROM component
  JOIN
  (
    SELECT cokey, sum((hzdepb_r - hzdept_r) * awc_r) AS component_whc,
    sum((hzdepb_r - hzdept_r)) AS depth
    FROM chorizon
    WHERE areasymbol = 'ca113'
    GROUP BY cokey
  ) AS a
  USING (cokey)
  WHERE component.areasymbol = 'ca113'
  -- filter out components that are missing soils data
  AND a.component_whc IS NOT NULL
  GROUP BY mukey
) AS b
-- JOIN constraint
USING (mukey)
-- optional constraint to limit geometry search in mapunit.poly table
WHERE mapunit.poly.areasymbol = 'ca113' ;

Create indexes and register the new geometry:
-- create attribute and spatial index:
CREATE UNIQUE INDEX yolo.whc_idx ON yolo.whc (ogc.fid) ;
CREATE INDEX yolo.whc_spatial_idx ON yolo.whc USING gist (wkb_geometry gist_geometry_ops);

-- register in geometry_columns table:
INSERT INTO geometry_columns VALUES ('','public','yolo.whc','wkb_geometry',2,9001,'POLYGON');
```

Soil Water Storage Computation



Additional Examples

- 1 simetaw work
- 2 soilweb
- 3 1km gridded soils data
- 4 PostGIS In Action Book
- 5 interactive examples
- 6 HRCLIM data: maybe for friday?